

HALO: Configuration-Calibrated Evidence Retrieval for Open-Label Human Activity Recognition

ABSTRACT

Human activity recognition systems are commonly tied to one sensor layout and one activity vocabulary. We study a deployment setting in which the device, placement, sampling regime, candidate activity names, and number of labeled demonstrations may all change. HALO represents acquisition configuration instead of normalizing it away. Its tokenizer computes frequency features on each stored sampling grid, defines bands in hertz, records acquisition-limited observability and patch-limited resolution, and folds the axes of each physical sensor into one token. We train the encoder without activity labels, using masked contextual prediction, masked sensor-description retrieval, and augmentation consistency.

Classification uses a versioned bank of sensor-patch embeddings. A closed-form evidence path first enforces modality and gravity compatibility, then uses a low-rank text-conditioned gate to decide whether a query and stored sensor can observe each candidate concept. It applies this admissibility test before cosine ranking, votes with enrolled-label identity or corpus-label text, and sums votes across query sensors. Adding a demonstration changes memory but does not update model parameters. The current Phase-A run uses 18 public datasets, 56 acquisition configurations, and 161 canonical labels; we reserve ten disjoint datasets for development and testing. On the three development datasets, the admissibility-gated path uses enrollment but trails both ungated retrieval and prototype or ridge classifiers on the same frozen features. We report this negative result rather than claim an unsupported retrieval advantage, and keep the seven-dataset test roster sealed until the method is fixed.

1 INTRODUCTION

A physiotherapist may ask a patient recovering from knee surgery to perform six exercises at home. The patient already owns a phone and a watch, but the exercise set was selected for that patient and may not appear in a public dataset. Collecting enough annotated examples to train a new model for every device and placement is rarely practical. Most inertial human activity recognition (HAR) systems assume the opposite setting: they fix the sensor layout and activity set, collect an annotated dataset, and fit a classifier for that combination.

We study a setting with three requirements. First, the activity set is *open*: an application supplies its candidate names at inference, including names absent from pretraining. Second, the input is *heterogeneous*: devices differ in placement, modality, channel availability, sampling regime, gravity convention, and filtering. Third, a deployment may provide zero, one, or a few labeled examples of each candidate. These requirements interact. A stored wrist example may be useful for recognizing an arm exercise from another wrist device, but it should not decide an arm gesture observed only by a pocket phone, however close their embeddings happen to be.

Two limitations motivate HALO.

(L1) Acquisition configuration changes what evidence is observable. Rotation, placement, sampling, firmware filtering, and sensor availability jointly change the recorded signal. A larger corpus improves coverage but cannot enumerate their continuous combinations. Nor does resampling make two instruments equivalent: interpolation cannot recover frequencies that the acquisition clock or device filter removed. We therefore expose selected configuration facts to both the representation and the retrieval rule.

(L2) A geometric neighbor is not automatically admissible evidence. Open-label methods often align sensor features with activity-name embeddings and classify by text similarity [46, 56, 84]. Yet language encodes objects and goals that an IMU may not observe, while small wording changes may distinguish a posture from a transition. We consequently compare recordings in sensor space, but condition the right to vote on whether the query and stored sensor configurations can observe the candidate concept. Section 2.3 gives concrete counterexamples without treating author-assigned motion relations as quantitative ground truth.

HALO. HALO has a label-free representation stage and an evidence stage. The tokenizer defines frequency bands in physical hertz on the stored grid, preserves signed low-frequency acceleration, and marks bands unavailable under the acquisition rate or unresolved by the patch duration. It folds the axes of each accelerometer or gyroscope into a sensor token with an explicit axis-validity mask. The token also receives a frozen text description and activity-invariant acquisition statistics.

Phase A learns from masked contextual prediction, reconstruction of hidden sensor descriptions, and consistency across physical augmentations.

At inference we freeze the encoder. A versioned memory stores one row per patch and sensor, with label and provenance. The evidence path applies hard modality/gravity compatibility, a learned low-rank admissibility gate, cosine ranking within the surviving rows, voting, and closed-form sensor fusion. An application appends a labeled demonstration to memory and binds it directly to its own candidate; no gradient update runs. Candidate names enter at run time rather than through a fixed classifier head.

We make four contributions.

- **We design a sensor-granular tokenizer** that defines frequency bands in physical hertz on each stored sampling grid, folds the axes of one physical sensor into a single token with an explicit validity mask, and marks the bands that the acquisition rate or the patch duration leaves unresolved. It states what a configuration *cannot* observe rather than presenting missing information as measured zero.
- **We separate admissibility from similarity** in the retrieval rule. A low-rank text-conditioned gate decides whether a query and a stored sensor can observe a candidate concept, and we apply that test *before* ranking so a rejected row frees its slot instead of merely losing weight.
- **We build a leakage-aware enrollment protocol** in which support and query come from different recorded executions, and we ship five controls with every number: ungated retrieval, support removal, support-label shuffling, prototypes, and ridge heads on the same frozen features.
- **We report a negative development result.** The fitted gate uses enrollment but does not beat those controls, and its telemetry shows the veto never fires. We bound the present claim, identify the protocol change that would make the mechanism testable, and keep the seven-dataset test roster sealed.

2 MOTIVATION

2.1 Application Requirements

Rehabilitation, occupational monitoring, and small field studies motivate an open activity set, several sensor configurations, and useful behavior from few demonstrations. The application, not the pretraining corpus, determines the relevant labels. Inertial traces are also difficult to

label in isolation: annotation commonly requires synchronized video, a scripted protocol, or self-report. A memory-based system can accept a small number of demonstrations without fitting a new model and can keep those recordings local to the deployment.

2.2 Acquisition Configuration Determines Observability

The recorded signal depends jointly on mounting orientation, body placement, sampling rate, device-side filtering, gravity convention, and sensor availability. Some differences are explicit in metadata; others are properties of the conversion. For example, one source may expose total acceleration while another exposes gravity-removed acceleration. Upsampling changes an array’s length but cannot restore bandwidth absent at acquisition. Likewise, full rotation invariance can discard posture information carried by gravity.

These facts motivate a distinction between *compatibility*, *admissibility*, and *similarity*. Compatibility asks whether two rows have comparable modalities and gravity states. Admissibility asks whether both sensor configurations can observe the candidate concept. A pocket accelerometer can be relevant to gait but may be uninformative for a fine arm gesture. Similarity ranks the remaining evidence. Applying these operations in the opposite order permits a geometrically close but physically irrelevant row to occupy a limited retrieval slot.

HALO therefore represents three kinds of acquisition information. Frequency bands are defined in hertz and carry masks for acquisition bandwidth and patch resolution. One frozen text descriptor states device, modality, placement, gravity convention, and partner-sensor presence. A separate activity-invariant vector measures acquisition behavior such as gravity magnitude, noise floor, quantization, clipping, rate fidelity, and rest bias. Missing measurements have explicit support bits. This split permits the paper to test whether the representation learned motion or merely recognized the source dataset.

2.3 Activity-Name Similarity Is Not a Motion Metric

An open-label model cannot rely on one fixed output head. A common alternative embeds candidate names and aligns sensor features with the same sentence space [46, 56, 84, 90]. That space reflects linguistic relation, not only observable motion. *Drinking coffee* and *drinking water* differ in object semantics while a wrist sensor may observe nearly the same gesture. Conversely, *sitting* and *sitting down* have similar words but denote a




Figure 1: Why name similarity is not used as a motion metric. These are qualitative counterexamples, not measured sensor distances. A quantitative geometry claim requires a predeclared all-pairs comparison between text embeddings and sensor-derived activity centroids; that artifact is not yet available.

posture and a transition. Figure 1 presents these as conceptual counterexamples; it does not assign quantitative motion distances by author judgment.

HALO uses language in two restricted roles: to describe acquisition configuration and to name candidates or corpus labels. It does not train the encoder to reproduce sentence distances. Stored recordings are ranked in sensor-feature space, while configuration-conditioned admissibility decides whether a neighbor may vote for a candidate. An enrolled demonstration is bound to the candidate by identity, so arbitrary names remain usable when labeled examples are present.

Dependency among components. Frequency coordinates and masks make rows from different sampling regimes interpretable. Sensor folding and descriptors expose configuration to the encoder. The evidence engine then uses that configuration to reject or downweight rows whose placement cannot observe a candidate. The thesis is therefore not that each component is individually novel, but that acquisition metadata calibrates evidence rather than being normalized away.

3 DESIGN OF HALO

3.1 System Overview

HALO separates representation learning from evidence-based classification (Figure 2). Phase A trains a sensor encoder without activity labels. The encoder is then frozen and used to build a versioned memory. At inference, a closed-form evidence path compares query and memory rows. The only fitted component in this path is a small admissibility gate that maps a sensor description and a candidate concept to one scalar. There is no relational decoder, learned retrieval projection, or deployment-specific output head in the active system.

For a window containing an accelerometer and gyroscope, the spectral front end first produces features for each scalar axis. The three axes of each modality are

Table 1: The active mechanisms and the distinction each preserves.

Problem	HALO mechanism
Different sampling regimes	Gaussian filterbank with logarithmic, constant- Q -inspired centers defined in hertz; acquisition observability and patch-resolution flags
Missing or damaged axes	Sensor folding with a three-bit axis-validity vector
Declared vs. measured configuration	Frozen sensor text plus separately measured, activity-invariant acquisition statistics
Irrelevant geometric neighbors	Compatibility and candidate-specific admissibility before top- k cosine ranking
Deployment-specific labels	Identity-bound enrolled votes and append-only memory

then folded into one token per patch and physical sensor. Frozen text describes what the sensor was declared to be, while a separate vector records activity-invariant acquisition measurements. A transformer contextualizes these tokens over physical time and across sensors. The memory stores the resulting sensor rows with labels and provenance.

3.2 Sensor-Granular Tokenizer

3.2.1 Frequency features in physical units. For each channel patch, the tokenizer removes the mean after applying a validity mask, applies a Hann window, and computes a zero-padded real DFT of size $S=256$. If the stored array rate is r , bin m represents mr/S hertz. It then evaluates $K=32$ Gaussian filters whose centers are logarithmically spaced from 0.3 to 15 Hz and whose nominal bandwidth scales with center frequency ($Q=4$). The construction is *constant- Q -inspired*: the cited constant- Q literature supports geometric center spacing [11, 12], while the Gaussian weighting and its parameters are HALO design choices.

The tokenizer operates on the stored grid and avoids an additional interpolation step in the model path. When a converter has already resampled a device stream, the stored rate still defines the DFT coordinate, but the original acquisition rate bounds observability. The representation therefore does not assert that different devices have identical filter responses; it only assigns the same physical coordinate to corresponding bands.

For center c_k and width $\sigma_k = c_k/(2Q)$, the observability mask is

$$o_k = \mathbf{1}[c_k + 2\sigma_k \leq 0.9r_{\text{source}}/2].$$

The patch-resolution flag is the clipped number of cycles in duration D , $q_k = \min(1, c_k D)$. This is a heuristic indication of spectral support, not a theorem that one cycle is sufficient. Unobservable bands and observed zero energy remain distinguishable.

[FIGURE PLACEHOLDER — TODO: redraw system overview]

Withdrawn 2026-08-17. The previous drawing depicted a two-resolution encoder and an active descriptor-prediction term; the current recipe uses a single 1s patch grid with that term disabled. Source preserved in `figures/fig_overview_v2artifact.tex`.

Figure 2: HALO end to end. Phase A learns one sensor representation without activity labels. A versioned bank stores frozen per-patch, per-sensor rows. Classification applies hard compatibility, candidate-conditioned admissibility, cosine ranking, voting, and sensor fusion. Enrollment appends labeled rows; it does not update parameters.

We retain the mean removed from the spectral path as a signed per-axis low-frequency feature for accelerometers. It can carry gravity and posture information. We do not attach the numerical posture-separation claims from an earlier draft because their generating artifact is absent; the ablation in Section 5.4 is the evidential test of this design choice.

3.2.2 Folding axes into physical sensors. Each accelerometer or gyroscope is a three-axis physical sensor. HALO concatenates the three per-axis patch tokens in fixed x, y, z order, zeros invalid axes, appends a three-bit validity vector, and projects the result to one token. Thus “axis reads zero” differs from “axis absent.” A six-channel accelerometer-gyroscope stream becomes two tokens per patch, while an accelerometer-only stream becomes one. Axis-role text is unnecessary after folding because axis identity is represented by its position inside the sensor token.

The augmentation stack respects this structure. We apply one $SO(3)$ rotation jointly to a triad and its co-located gyroscope, and we permit sensor masking only when another modality at the same placement remains visible, so the model never has to reconstruct a wrist from an ankle or a pocket.

3.2.3 Declared descriptors and measured acquisition bias. Each sensor receives two frozen artifacts. A 384-dimensional sentence embedding describes device, modality, placement, gravity state, and whether a partner modality was present. A 14-dimensional measurement vector contains seven standardized statistics and seven support bits: gravity magnitude, gravity presence, high-band noise floor, quantization step, clipping fraction, native-to-stored rate fidelity, and rest bias. These statistics are estimated on training subjects and, where possible, quiescent windows; unsupported fields remain unsupported rather than being imputed as measurements.

Separate learned MLPs project the two artifacts into the token width and inject them through gated residuals. Keeping the artifacts frozen preserves their audit

meaning. The representation may still encode dataset provenance, so we measure dataset-identification accuracy and keep similarity in the acquisition-statistics vector disabled in the active retriever unless a provenance guard shows it is safe.

3.2.4 Context encoder. The transformer encoder [73] has width 256, six layers, eight heads, and a frozen all-MiniLM-L6-v2 text encoder [57]. Attention alternates over time within a sensor and across sensors at a patch; padding and sensor masks exclude missing values. Rotary positions use elapsed seconds rather than token indices [70]. Each window is represented at two patch-duration resolutions. Pooling weights real patch duration within a resolution and gives equal weight to active resolutions.

3.3 Phase A: Label-Free Pretraining

Phase A has two top-level objectives. Their coefficients are calibrated once at step 2,000 from encoder-gradient norms and then fixed.

Masked contextual prediction. A student predicts contextual sensor tokens produced by an exponential-moving-average teacher [2, 3, 26]. Masking includes contiguous time blocks and, when physically well posed, a complete modality at one placement. A descriptor-mask event hides a sensor’s text while leaving its signal and measured bias visible; an InfoNCE retrieval loss [72] then ranks the true frozen descriptor among the distinct descriptions in the batch. This term is part of the JEPa objective family rather than a third top-level objective.

Augmentation consistency. VICReg [6] aligns two physical augmentations of the same window while maintaining feature variance and reducing covariance. It avoids declaring other windows to be negatives when they may contain the same activity. Activity labels are not used in either gradient path. Checkpoint selection uses subject-disjoint k -nearest-neighbor balanced accuracy on the Phase-A validation subjects.




Figure 3: Active evidence path. Compatibility and admissibility are applied before candidate-specific top- k ranking. Enrolled rows vote only for the candidate to which the application bound them; corpus rows vote through label-text similarity. Per-sensor votes are summed.

3.4 Admissibility-Gated Evidence

Bank. The current schema stores one row per patch and sensor: a 256-dimensional feature, the frozen sensor descriptor, acquisition measurements, label, subject, dataset, configuration, recording event, time, and resolution. We bind the bank to the exact encoder checkpoint and schema, and exclude query-window, event, and subject identifiers wherever the evaluation protocol demands it. Enrollment adds rows whose episode-local candidate identity is explicit.

Train-only resolvability measurement. For every training sensor and concept with adequate support, a subject-disjoint one-versus-rest k NN probe measures how well that sensor distinguishes the concept. Chance balanced accuracy maps to zero. We measure accelerometers and gyroscopes separately and report simultaneous-placement contrasts separately. Activity strings only group the measurement cells; this probe never embeds them.

Admissibility gate. A rank- r bilinear model maps frozen sensor and concept text to one weight:

$$u = Ae_s, \quad v = Be_c, \quad g(s, c) = \sigma(u^\top \Lambda v + b),$$

where $e_s, e_c \in \mathbb{R}^{384}$ and $r=2$ in the active artifact. The low rank limits the gate’s ability to memorize literal stream-label cells. We fit it only to the train-only resolvability table, and for text unlike the fitted sensor or concept supports we blend the prediction toward the table’s neutral median rather than extrapolating confidently. Characterizing generalization requires held-out configuration, dataset, and concept folds; the fitted table itself is not validation evidence.

Prediction. For each query sensor, HALO performs the following operations in order.

- (1) Hard compatibility removes rows with incompatible modality or gravity state. This step is label independent.
- (2) The gate scores the query sensor and memory sensor against each candidate. Their geometric mean requires both sides to observe the concept. Rows below the veto threshold are removed *before* top- k , so a rejected row frees its slot.
- (3) Feature cosine ranks the surviving rows separately for each candidate. The active budget is 64 rows per query patch and candidate. Similarity in the acquisition-statistics vector has weight zero in the reported artifact.
- (4) An enrolled row votes for its bound candidate by identity. A corpus row votes through the rectified cosine between its stored label text and the candidate text; this is the only $k=0$ bridge.
- (5) Votes are summed across the query’s sensors. This fusion is closed form and does not infer one placement’s signal from another.

Setting every admissibility value to one gives the ungated identity-retrieval control. Prototype and ridge classifiers use the same frozen representations and enrolled examples. Support removal and cyclic support-label shuffling test whether an apparent enrollment gain actually uses the labeled rows.

Scope. The active system does not include an unknown output, confidence calibration, Core ML export, an iOS application, or a public `predict()` API. Selective classification would require truth-absent episodes and risk-coverage evaluation [23]; evidential deep learning is a distinct uncertainty method [60] and is not evidence for the present retrieval rule.

4 IMPLEMENTATION

We implement HALO in PyTorch as a single repository covering conversion, pretraining, memory construction, and the evidence rule. The system is deliberately small where it matters: the trunk holds 7.70 M parameters and the only component fitted after Phase A is a rank-2 bilinear gate with roughly 1.5 k parameters.

4.1 Codebase

Table 2 lists the components. We write one converter per dataset — 35 in total — rather than a generic ingest path, because every public corpus differs in column order, unit convention, gravity handling, and subject encoding, and a generic reader hides exactly the errors that later corrupt a physical feature. Each converter emits whole-recording sessions; the grid builder owns all windowing, so no converter can silently choose its own window policy.

Table 2: Implementation components. Line counts exclude generated artifacts.

Component	Lines	Contents
Dataset converters	11.8 k	35 per-dataset readers, unit and gravity normalization
Corpus pipeline	6.9 k	curation policy, windowing, three grid alignments, quality screens
Model	3.9 k	filterbank, sensor folding, conditioning, dual-branch trunk
Phase-A training	7.6 k	sampler, augmentation, masking, objectives, telemetry
Evidence stage	14.3 k	bank build, resolvability, gate, prediction rule, protocol
Evaluation harness	2.1 k	baseline adapters, scoring, protocol fingerprinting
Tests	9.3 k	578 tests across 45 files

4.2 Training and Memory Cost

We train on one NVIDIA RTX 4090 (24 GB). The reported Phase-A run completes 30,000 steps at batch 256 in 33 minutes of wall clock and peaks at 4.36 GiB of allocated device memory, so a full pretraining run fits comfortably on a single consumer GPU. Mixed-precision training is enabled and we compile only the transformer; compiling the encoder specializes on each batch’s changing unique-text cardinality and runs slower.

Memory construction is the expensive artifact. The schema-4 bank stores one row per patch and sensor, which is 7.65 M rows for a 250,000-window archive and 6.9 GB on disk in half precision. Two engineering choices keep the evidence path tractable. First, we do not store a 384-dimensional sensor descriptor on every row: roughly 10^5 rows are drawn from about 40 distinct sensor descriptions, so we store a de-duplicated table and an index, evaluate the gate once per distinct description, and gather. This is exact, not approximate, because the gate reads only text. Second, we cache stream metadata and text encodings per configuration, so sentence-embedding inference never runs inside a query loop.

4.3 Artifact Binding

A retrieval system built on a cached bank fails in a specific way: an encoder is retrained, the bank is not rebuilt, and every downstream number silently describes a stale embedding space. We make that failure impossible to reach rather than easy to notice.

Every artifact carries a fingerprint of what produced it, and every consumer checks it before doing any work. The bank records the encoder checkpoint hash, the label vocabulary, and behavioral probes of both the embedding path and the sentence-embedding path; loading it re-runs those probes against the live encoder and refuses a mismatch. The gate records the resolvability table hash, the checkpoint that produced it, and the bank it is

bound to. The evaluation protocol is itself fingerprinted over the source files that define scoring behavior, so a change to the protocol invalidates cached results instead of blending two protocols into one table. The quality screens are loaded with the same discipline: a missing or stale exclusion cache is a hard failure, because silently re-admitting byte-identical stale-buffer windows would corrupt both training and scoring.

These guards are load-bearing rather than decorative. During development they rejected a legacy resolvability table that pooled evaluation datasets into a training-only measurement, and they currently refuse the gate artifact described in Section 5.3 whenever it is paired with a bank from a different checkpoint.

4.4 Reproducibility

Every quantitative statement in this paper is bound to a named artifact through a checked-in evidence ledger that records the producing command, the source-data manifest, the seed, the checkpoint fingerprint, and the raw output. Values without that chain are not printed as results, which is why several quantities from an earlier draft of this work do not appear here. We will release the converters, the corpus pipeline, the model, the evidence stage, and the evaluation harness.

5 EVALUATION

We ask a narrower question than a completed benchmark paper answers: does the active admissibility-gated rule improve on simpler enrollment rules on the development roster? We report that result here. We leave the external test roster and several tokenizer ablations unevaluated, and we state them as pending measurements rather than filling the space with fabricated cells.

5.1 Data and Artifacts

Phase A uses 18 public datasets and 56 acquisition configurations. Their materialized native-grid manifests contain 1,871,011 stream-window rows before quality screens and the subject split. The selected bank contains 250,000 balanced source windows, 546 training subjects, all 161 canonical training labels, 3,942,901 patch rows, and 7,653,242 sensor rows. Ten datasets are disjoint from Phase-A training: three are the development roster and seven are sealed for testing. The full grid-derived roster appears in Table 3; its machine-readable provenance is listed in the paper-source evidence ledger.

The completed run trains for 30,000 steps. The subject-disjoint Phase-A selection metric peaks at step 4,000 (label-stream-macro balanced accuracy 0.3028),

Table 3: Current Phase-A and external-evaluation rosters. *Cfg* counts materialized sensor configurations; *rows* counts stream-window rows before the current quality exclusions and subject split, so simultaneous placements are counted separately. Counts describe the grid build the reported checkpoint was trained on. [TODO: regenerate with the checkpoint repin.] [†]The XRF ear stream was acquired at 25 Hz and stored at 50 Hz. [‡]MM-Fit devices have different acquisition clocks and are placed on a shared 100 Hz grid; source rates still bound observability. The seven test datasets have not been used for method selection.

Phase-A training				External evaluation			
Dataset	Hz	Cfg	rows	Dataset (role)	Hz	Cfg	rows
CAPTURE-24 [15]	100	1	1,530,792	MotionSense [43] (dev)	50	1	4,534
WISDM [77]	20	2	70,205	RealWorld [71] (dev)	50	1	11,237
RealDISP [4]	50	9	50,886	Shoaib [62] (dev)	50	1	2,100
DSADS [7]	25	5	37,815	InclusiveHAR [30] (test)	50	1	1,260
XRF V2 [36]	25/50 [†]	6	34,764	USC-HAD [88] (test)	100	1	4,350
NFI-FARED [45]	100	2	26,876	TNDA-HAR [85] (test)	50	1	3,353
PhyTMO [22]	100	8	24,684	UT-Complex [63] (test)	50	1	3,900
HARMES [13]	50	1	18,576	MoniPar [64] (test)	50	1	12,079
UniMiB SHAR [47]	50	1	11,771	SPAR [14] (test)	50	2	1,898
Opportunity [16]	30	5	10,765	Upper Limb Use [18] (test)	50	4	1,124
FORTH-TRACE [32]	51.2	5	10,340				
UCI-HAR [1]	50	1	10,299				
HHAR [68]	50	1	9,587				
KU-HAR [65]	100	1	9,324				
MM-Fit [69]	100 [‡]	4	6,280				
SP-SW-HAR [44]	100	2	3,757				
PAMAP2 [58]	100	1	3,186				
MHEALTH [5]	50	1	1,104				

so `best.pt` rather than the last checkpoint builds the reported bank. Dense development-only comparisons also favor this checkpoint; the test roster is not involved in that choice.

5.2 Evaluation Protocol

An evaluation dataset supplies its own candidate inventory; this is a HALO protocol choice, not a claim established by the generalized zero-shot-learning benchmark. Following that benchmark, we separately report whether a candidate was seen or unseen during training [78]. Macro-F1 is primary because activity frequencies are imbalanced.

For enrollment, support and query executions are disjoint. The development evaluation uses cross-subject, same-configuration cohorts for MotionSense, RealWorld, and Shoaib. At each $k \in \{1, 2, 4, 8\}$, we retain the same queries and candidate inventory and nest the support prefixes. A verified recording execution, not an adjacent window, counts as one example. We also run support removal, cyclic support-label shuffling, an all-one admissibility control, prototypes, and ridge heads on the identical features and support split.

The test roster was fixed before the present gate evaluation: InclusiveHAR, USC-HAD, TNDA-HAR, UT-Complex, MoniPar, SPAR, and Upper Limb Use. TNDA-HAR has no usable subject-disjoint enrollment cohort and contributes only to a future zero-support cell.

Table 4: Development macro-F1 (percentage points), averaged over MotionSense, RealWorld, and Shoaib. These are model-selection diagnostics, not sealed-test results.

Method	labeled executions per candidate (k)				
	0	1	2	4	8
HALO admissibility gate	15.73	40.88	43.11	44.29	39.72
HALO ungated retrieval	9.07	52.15	57.14	56.16	53.62
Frozen HALO + prototype	—	57.76	61.16	65.41	72.04
Frozen HALO + ridge	—	56.62	59.42	62.60	71.29
Gate, support removed	—	15.73	15.73	15.73	15.73
Gate, support labels shuffled	—	6.02	4.75	2.19	0.88

SPAR’s repetitions occur within one bout and do not establish across-session enrollment. MoniPar is the held-out source that can support across-week enrollment. We do not consume any of these test datasets in the present revision.

5.3 Development Result: the Gate Does Not Beat Simpler Controls

Table 4 averages dataset-level macro-F1 percentages over the three development datasets. “Ungated” uses the identical prediction path with admissibility fixed to one. Prototype and ridge operate on the same frozen representation and enrolled executions.

The support interventions show that the gated rule uses enrollment: removing or shuffling support substantially reduces macro-F1. That fact is insufficient for a performance claim. Ungated retrieval is 11.3–14.0 points better at $k = 1$ through 4, and prototypes are 16.9–32.3 points better across the enrollment curve. Performance also falls from $k = 4$ to 8 for the gated rule while prototype and ridge improve.

The telemetry identifies a concrete mechanism. Across every reported development cell, gate values range from 0.245 to 0.733, above the configured veto threshold 0.15; consequently the hard veto removes no row. The gate only reweights already selected evidence and cannot demonstrate the paper’s intended admissibility behavior. We therefore reject the hypothesis that the current gate improves few-example recognition. The seven-dataset test evaluation remains sealed rather than being spent on a method that fails its development controls.

5.4 Required Measurements Before a Performance Claim

The following experiments are part of the protocol but do not yet have provenance-complete artifacts:

- held-out concept, dataset, and body-region studies for the gate, including gate-collapse, provenance-capture, and latent-to-measurement guards;
- cross-configuration enrollment on simultaneous training streams, reported separately from the external test result;
- native-grid versus common-rate, sensor-folded versus channel-level, descriptor-off, acquisition-bias-off, observability-off, and signed-low-frequency ablations;
- all-pairs activity-text similarity versus sensor-centroid similarity, replacing any hand-selected quantitative geometry claim;
- data-matched and released-checkpoint baselines, including the published short fine-tuning baseline required by the wearable-foundation-model literature; and
- only after development criteria are met, one evaluation of the seven-dataset sealed test roster with confidence intervals and seen/unseen candidate breakdowns.

No on-device latency, Core ML, iOS, or unknown-class result is claimed because the corresponding implementation and measurement artifacts do not exist in the repository.

6 RELATED WORK

Cross-dataset and foundation models for wearable sensing. DeepConvLSTM [53] is representative of supervised HAR models trained for one benchmark. Later work addresses selected sources of variation [54, 55, 79] or pretrains transferable encoders. LiMU-BERT [80] uses masked reconstruction at 20 Hz, CrossHAR [27] adds a contrastive objective, and ssl-wearables [86] scales self-supervision on wrist accelerometry. Scaling Wearable Models [49] studies population-scale pretraining and imputation; LSM-2 [82] specifically targets incomplete wearable observations. RelCon [81] learns relative motion similarity from roughly one billion sensor segments, and SensorLM [91] pairs wearable signals with language. Tartan IMU [92] studies inertial positioning, while general time-series architectures include MOMENT and PatchTST [24, 51]. Mitra is not included in this group because it is a tabular, rather than time-series, foundation model [89].

The controlled Inertia-1 study varies modality, placement, rate, and window length [83]; the recent HAR foundation-model survey provides a broader taxonomy [8]. We cite the study for its controlled findings and the survey only for the landscape. HALO’s narrower question is whether acquisition configuration can calibrate which retrieved evidence is allowed to vote.

Sampling and acquisition context. Rate-invariant autoencoding disentangles order-preserving temporal warp into rate parameters and a rate-invariant latent [34]. FlowState combines a state-space model with a functional basis decoder and dynamic timescale adjustment for sampling-rate-equivariant forecasting [25]. Universal Spectral Tokenization accepts native wavelength grids in astronomy and proposes extension to other sampled domains; it is adjacent inspiration, not evidence on wearable HAR [61]. MoPFormer learns IMU motion primitives after resampling its datasets to 100 Hz [87]. These systems therefore should not be summarized as sharing one rate-normalization strategy.

Other work supplies acquisition context explicitly. GOAT [46] uses activity and body-position text, ZeroHAR [17] uses categorical sensor context, and CHARM was introduced as a preprint and subsequently presented in the peer-reviewed SensorsVoice work [19, 20]. SensorLLM [39] and oneHAR [76] also use language for sensor context, whereas UniMTS [90] represents placement as a graph and uses rotation augmentation. HALO does not claim that sensor descriptions are new. Its contribution is to separate a declared description from measured acquisition statistics and then use configuration to calibrate retrieval admissibility.

Retrieval and label-efficient recognition. Template matching has long been used for gestures [29, 40, 50]. Matching Networks and Prototypical Networks learn or use comparison rules under episodic tasks [67, 75]. Language and vision systems augment parametric models with external non-parametric memory [10, 28, 33, 37, 48]; not all of these systems keep the parametric model frozen. Memory-modular classification studies adding classes by changing memory contents [31].

In sensing, RAG-HAR retrieves signal descriptors for an LLM [66], and ZARA is a training-free, evidence-grounded LLM agent [38]. TimEE and TIC-FM perform zero- or few-example time-series classification in context [21, 35]. These are not evidence for a target-example k curve without parameter updates unless such a curve is explicitly reported. Appendix A therefore separates open-label evaluation, use of target examples, and no-update enrollment instead of collapsing them into one checkmark.

7 DISCUSSION AND LIMITATIONS

What the development result rules out, and what it does not. The gated rule uses enrollment — removing or shuffling support collapses it — but it trails ungated retrieval, prototypes, and ridge heads on the same frozen features. We therefore reject the claim that the current gate improves few-example recognition. We do not conclude that admissibility is the wrong idea, because the protocol that produced this result cannot express it. Every scored cell enrolls and queries the *same* configuration, so the evidence row carries the query’s own descriptor and the two-sided gate collapses to a per-candidate constant. A query-independent constant multiplying an evidence-derived vote can only add bias. The cross-configuration cell that would separate these two readings is pure evaluation: it needs no retraining, only support drawn from a different device than the query.

The veto never fires, and the threshold is not calibrated to the artifact. Gate values on the development cells span 0.245 to 0.733 against a threshold of 0.15, so no row is ever removed. The mechanism reduces to reweighting. Worse, the threshold is a fixed constant while the fitted gate’s scale moves with the encoder that produced the resolvability table: on a later table the fitted neutral value sits at 0.190, only marginally above the same threshold, and the fraction of fitted cells below it doubles. Nothing in the current design ties the veto to the distribution it gates. A threshold expressed as a quantile of the fitted values, with an asserted ordering between the neutral fallback and the veto, would remove

a failure mode in which an unfamiliar configuration is either always admitted or vetoed on every candidate.

Provenance is confounded with hardware in public corpora. Each public dataset generally uses one device model and one preprocessing chain, so a descriptor of acquisition behavior is also a descriptor of which study a recording came from. Our encoder reads dataset identity from pooled embeddings well above chance. We respond by keeping the acquisition-statistics similarity term disabled in the reported retriever and by measuring same-dataset retrieval directly rather than assuming it is benign, but the confound is a property of the available data and cannot be resolved by a better estimator. Settling it needs a corpus with the same hardware across studies, or the same study across hardware.

Scope of the recognition problem. HALO classifies presegmented single-activity windows. It has no unknown output and does not perform continuous segmentation or calibrated rejection, so a deployment that encounters an activity outside its candidate set will still name one. Selective classification would require truth-absent episodes and risk-coverage evaluation [23]; we treat it as future work rather than claiming it, because the evaluation artifacts that would support such a claim do not exist.

No on-device claim. We report training and memory cost on a workstation GPU. We make no claim about phone-side latency, energy, or export, because we have not built that path. The design does bound what such a claim would look like: the encoder is 7.70 M parameters and enrollment appends rows rather than running a gradient step, so the runtime question is retrieval over a bank rather than on-device training. The bank is also the obvious pressure point — 6.9 GB is a server-side artifact, and a phone-side deployment would need either a per-user subset or a compressed index.

Sealing the test roster. Seven evaluation datasets remain unconsumed. We chose to report a negative development result rather than spend them on a method that fails its own controls, and we state the criteria for unsealing in Section 5.4. This is a deliberate cost: it means this paper reports no external performance number. We think that is the correct trade when the alternative is a headline produced by a mechanism whose telemetry shows it is inactive.

8 CONCLUSION

HALO tests one proposition: acquisition configuration should calibrate evidence rather than be normalized away. We build a sensor-granular encoder that defines

frequency bands in physical hertz, marks what a configuration cannot resolve, and learns without activity labels; we classify by a transparent sequence of compatibility, admissibility, cosine ranking, voting, and sensor fusion; and we enroll a new activity by appending rows to memory rather than updating parameters.

The proposition remains open. Our retrieval stage works — it uses enrollment, and its ungated form beats prototype and ridge heads on the same frozen features — but the admissibility gate that carries the thesis does not beat those controls, and its telemetry shows the veto never fires. We report that rather than a headline, and we keep seven evaluation datasets sealed. The next measurement is cheap and decisive: enroll on one device and query from another, which is the only regime in which a two-sided admissibility test can express anything at all.

REFERENCES

- [1] Davide Anguita, Alessandro Ghio, Luca Oneto, Xavier Parra, and Jorge Luis Reyes-Ortiz. 2013. A Public Domain Dataset for Human Activity Recognition Using Smartphones. In *Proceedings of the European Symposium on Artificial Neural Networks, Computational Intelligence and Machine Learning (ESANN)*. Bruges, Belgium, 437–442.
- [2] Mahmoud Assran, Quentin Duval, Ishan Misra, Piotr Bojanowski, Pascal Vincent, Michael Rabbat, Yann LeCun, and Nicolas Ballas. 2023. Self-Supervised Learning from Images with a Joint-Embedding Predictive Architecture. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*. IEEE, Vancouver, Canada, 15619–15629. <https://doi.org/10.1109/CVPR52729.2023.01499>
- [3] Alexei Baevski, Wei-Ning Hsu, Qiantong Xu, Arun Babu, Jiatao Gu, and Michael Auli. 2022. data2vec: A General Framework for Self-supervised Learning in Speech, Vision and Language. In *Proceedings of the 39th International Conference on Machine Learning (ICML)*. PMLR, Baltimore, MD, 1298–1312.
- [4] Oresti Baños, Miguel Damas, Héctor Pomares, Ignacio Rojas, Máté Attila Tóth, and Oliver Amft. 2012. A Benchmark Dataset to Evaluate Sensor Displacement in Activity Recognition. In *Proceedings of the 2012 ACM Conference on Ubiquitous Computing*. ACM, 1026–1035. <https://doi.org/10.1145/2370216.2370437>
- [5] Oresti Banos, Rafael Garcia, Juan A. Holgado-Terriza, Miguel Damas, Hector Pomares, Ignacio Rojas, Alejandro Saez, and Claudia Villalonga. 2014. mHealthDroid: A Novel Framework for Agile Development of Mobile Health Applications. In *Proceedings of the 6th International Work-conference on Ambient Assisted Living and Daily Activities (IWAAL) (LNCS)*, Vol. 8868. Springer, Belfast, UK, 91–98. https://doi.org/10.1007/978-3-319-13105-4_14
- [6] Adrien Bardes, Jean Ponce, and Yann LeCun. 2022. VICReg: Variance-Invariance-Covariance Regularization for Self-Supervised Learning. In *The Tenth International Conference on Learning Representations (ICLR)*. OpenReview.net, Virtual.
- [7] Billur Barshan and Murat Cihan Yükses. 2014. Recognizing Daily and Sports Activities in Two Open Source Machine Learning Environments Using Body-Worn Sensor Units. *Comput. J.* 57, 11 (2014), 1649–1667. <https://doi.org/10.1093/comjnl/bxt075>
- [8] Sizhen Bian, Mengxi Liu, Lala Shakti Swarup Ray, Bo Zhou, Bin Guo, Zhiwen Yu, Thomas Plötz, Paul Lukowicz, Siyu Yuan, and Vitor Fortes Rey. 2026. Foundation Models Defining a New Era in Sensor-based Human Activity Recognition: A Survey and Outlook. *Proceedings of the ACM on Interactive, Mobile, Wearable and Ubiquitous Technologies (IMWUT)* 10, 2, Article 31 (2026), 38 pages. <https://doi.org/10.1145/3810230> arXiv:2604.02711.
- [9] Sizhen Bian and Paul Lukowicz. 2025. RecGym: Gym Workouts Recognition Dataset with IMU and Capacitive Sensor. UCI Machine Learning Repository. (2025). <https://doi.org/10.24432/C5PW4K> <https://archive.ics.uci.edu/dataset/1128>
- [10] Sebastian Borgeaud, Arthur Mensch, Jordan Hoffmann, Trevor Cai, Eliza Rutherford, Katie Millican, George van den Driessche, Jean-Baptiste Lespiau, Bogdan Damoc, Aidan Clark, Diego de Las Casas, Aurelia Guy, Jacob Menick, Roman Ring, Tom Hennigan, Saffron Huang, Loren Maggiore, Chris Jones, Albin Cassirer, Andy Brock, Michela Paganini, Geoffrey Irving, Oriol Vinyals, Simon Osindero, Karen Simonyan, Jack W. Rae, Erich Elsen, and Laurent Sifre. 2022. Improving Language Models by Retrieving from Trillions of Tokens. In *Proceedings of the 39th International Conference on Machine Learning (ICML)*. PMLR, Baltimore, MD, 2206–2240.
- [11] Judith C. Brown. 1991. Calculation of a Constant Q Spectral Transform. *The Journal of the Acoustical Society of America* 89, 1 (1991), 425–434. <https://doi.org/10.1121/1.400476>
- [12] Judith C. Brown and Miller S. Puckette. 1992. An Efficient Algorithm for the Calculation of a Constant Q Transform. *The Journal of the Acoustical Society of America* 92, 5 (1992), 2698–2701. <https://doi.org/10.1121/1.404385>
- [13] Robin Burchard, Pascal-André Brückner, Marius Bock, Jürgen Gall, and Kristof Van Laerhoven. 2026. HARMES: A Multi-Modal Dataset for Wearable Human Activity Recognition with Motion, Environmental Sensing and Sound. *arXiv preprint arXiv:2605.02596* (2026).
- [14] David M. Burns, Nathan Leung, Michael Hardisty, Cari M. Whyne, Patrick Henry, and Stewart McLachlin. 2018. Shoulder Physiotherapy Exercise Recognition: Machine Learning the Inertial Signals from a Smartwatch. *Physiological Measurement* 39, 7 (2018), 075007. <https://doi.org/10.1088/1361-6579/aacfd9>
- [15] Shing Chan, Hang Yuan, Catherine Tong, Aidan Acquah, Abram Schonfeldt, Jonathan Gershuny, and Aiden Doherty. 2024. CAPTURE-24: A large dataset of wrist-worn activity tracker data collected in the wild for human activity recognition. *Scientific Data* 11, 1 (2024), 1135. <https://doi.org/10.1038/s41597-024-03960-3>
- [16] Ricardo Chavarriaga, Hesam Sagha, Alberto Calatroni, Sundara Tejaswi Digumarti, Gerhard Tröster, José del R. Millán, and Daniel Roggen. 2013. The Opportunity Challenge: A Benchmark Database for On-Body Sensor-Based Activity Recognition. *Pattern Recognition Letters* 34, 15 (2013), 2033–2042. <https://doi.org/10.1016/j.patrec.2012.12.014>
- [17] Ranak Roy Chowdhury, Ritvik Kapila, Ameya Panse, Xiyuan Zhang, Diyan Teng, Rashmi Kulkarni, Dezhi Hong, Rajesh K. Gupta, and Jingbo Shang. 2025. ZeroHAR: Sensor Context Augments Zero-Shot Wearable Action Recognition. In

- Proceedings of the 39th AAAI Conference on Artificial Intelligence (AAAI)*, Vol. 39. AAAI Press, Philadelphia, PA, 16046–16054. <https://doi.org/10.1609/aaai.v39i15.33762>
- [18] Ann David, Stephen Sukumaran Reetha Janet Sureka, Sankaralingam Gayathri, Salai Jeyseelan Annamalai, Selvaraj Samuel Kamlesh Kumar, Anju Kuruvilla, Henry Prakash Magimairaj, S. K. M. Varadhan, and Sivakumar Balasubramanian. 2021. Quantification of the Relative Arm Use in Patients with Hemiparesis Using Inertial Measurement Units. *Journal of Rehabilitation and Assistive Technologies Engineering* 8 (2021), 20556683211019694. <https://doi.org/10.1177/20556683211019694>
 - [19] Utsav Dutta, Sina Khoshfetrat Pakazad, and Henrik Ohlsson. 2025. Time to Embed: Unlocking Foundation Models for Time Series with Channel Descriptions. *arXiv preprint arXiv:2505.14543* (2025). <https://doi.org/10.48550/arXiv.2505.14543> System name: CHARM.
 - [20] Utsav Dutta, Gerardo Pastrana, Sina Khoshfetrat Pakazad, and Henrik Ohlsson. 2026. Giving Sensors a Voice: Multimodal JEPA for Semantic Time-Series Embeddings. In *Proceedings of the 43rd International Conference on Machine Learning (ICML) (Proceedings of Machine Learning Research)*, Vol. 306. PMLR, Seoul, South Korea. arXiv:2605.31580.
 - [21] Juntao Fang, Shifeng Xie, Shengbin Nie, Yuhui Ling, Yuming Liu, Zijian Li, Keli Zhang, Lujia Pan, Themis Palpanas, and Ruichu Cai. 2026. Rethinking Zero-Shot Time Series Classification: From Task-specific Classifiers to In-Context Inference. *arXiv preprint arXiv:2602.00620* (2026). System name: TIC-FM.
 - [22] Sara García-de Villa, Ana Jiménez-Martín, and Juan Jesús García-Domínguez. 2022. A Database of Physical Therapy Exercises with Variability of Execution Collected by Wearable Sensors. *Scientific Data* 9, 1 (2022), 266. <https://doi.org/10.1038/s41597-022-01387-2>
 - [23] Yonatan Geifman and Ran El-Yaniv. 2017. Selective Classification for Deep Neural Networks. In *Advances in Neural Information Processing Systems 30 (NIPS)*. Curran Associates, Long Beach, CA, 4878–4887.
 - [24] Mononito Goswami, Konrad Szafer, Arjun Choudhry, Yifu Cai, Shuo Li, and Artur Dubrawski. 2024. MOMENT: A Family of Open Time-Series Foundation Models. In *Proceedings of the 41st International Conference on Machine Learning (ICML)*. PMLR, Vienna, Austria, 16115–16152.
 - [25] Lars Graf, Thomas Ortner, Stanisław Woźniak, and Angeliki Pantazi. 2026. FlowState: Sampling-Rate-Equivariant Time-Series Forecasting. In *Proceedings of the 43rd International Conference on Machine Learning (ICML) (Proceedings of Machine Learning Research)*, Vol. 306. PMLR, Seoul, South Korea. arXiv:2508.05287.
 - [26] Jean-Bastien Grill, Florian Strub, Florent Altché, Corentin Tallec, Pierre H. Richemond, Elena Buchatskaya, Carl Doersch, Bernardo Ávila Pires, Zhaohan Guo, Mohammad Gheshlaghi Azar, Bilal Piot, Koray Kavukcuoglu, Rémi Munos, and Michal Valko. 2020. Bootstrap Your Own Latent: A New Approach to Self-Supervised Learning. In *Advances in Neural Information Processing Systems 33 (NeurIPS)*. Curran Associates, Virtual, 21271–21284.
 - [27] Zhiqing Hong, Zelong Li, Shuxin Zhong, Wenjun Lyu, Haotian Wang, Yi Ding, Tian He, and Desheng Zhang. 2024. CrossHAR: Generalizing Cross-dataset Human Activity Recognition via Hierarchical Self-Supervised Pretraining. *Proceedings of the ACM on Interactive, Mobile, Wearable and Ubiquitous Technologies (IMWUT)* 8, 2, Article 64 (2024), 26 pages. <https://doi.org/10.1145/3659597>
 - [28] Ahmet Iscen, Mathilde Caron, Alireza Fathi, and Cordelia Schmid. 2024. Retrieval-Enhanced Contrastive Vision-Text Models. In *The Twelfth International Conference on Learning Representations (ICLR)*. OpenReview.net, Vienna, Austria.
 - [29] Holger Junker, Oliver Amft, Paul Lukowicz, and Gerhard Tröster. 2008. Gesture spotting with body-worn inertial sensors to detect user activities. *Pattern Recognition* 41, 6 (2008), 2010–2024. <https://doi.org/10.1016/j.patcog.2007.11.016>
 - [30] Seyed Reza Kamel Tabbakh Farizani, Iman Naeimi, Kosar Naghavi, Fatemeh Majidi Nasab, Mehran Ghaffarian, and Mahdih Sadat Sharifi Moghaddam Kakhki. 2026. InclusiveHAR: A Smartphone-Based Dataset for Human Activity Recognition Across Diverse Physical Abilities. Mendeley Data, V4. (2026). <https://doi.org/10.17632/r78dn3f6nc.4>
 - [31] Dahyun Kang, Ahmet Iscen, Eunchan Jo, Sua Choi, Minsu Cho, and Cordelia Schmid. 2025. Memory-Modular Classification: Learning to Generalize with Memory Replacement. *Transactions on Machine Learning Research (TMLR)* (2025). arXiv:2504.06021.
 - [32] Katerina Karagiannaki, Athanasia Panousopoulou, and Panagiotis Tsakalides. 2016. A Benchmark Study on Feature Selection for Human Activity Recognition. In *Proceedings of the 2016 ACM International Joint Conference on Pervasive and Ubiquitous Computing: Adjunct*. ACM, 105–108. <https://doi.org/10.1145/2968219.2971421>
 - [33] Urvashi Khandelwal, Omer Levy, Dan Jurafsky, Luke Zettlemoyer, and Mike Lewis. 2020. Generalization through Memorization: Nearest Neighbor Language Models. In *The Eighth International Conference on Learning Representations (ICLR)*. OpenReview.net, Addis Ababa, Ethiopia.
 - [34] Kaushik Koneripalli, Suhas Lohit, Rushil Anirudh, and Pavan Turaga. 2020. Rate-Invariant Autoencoding of Time-Series. In *Proceedings of the IEEE International Conference on Acoustics, Speech and Signal Processing (ICASSP)*. IEEE, Barcelona, Spain, 3732–3736. <https://doi.org/10.1109/ICASSP40776.2020.9053983>
 - [35] Jaris Küken, Shi Bin Hoo, Martin Mráz, Frank Hutter, and Lennart Purucker. 2026. TimEE: End-to-end Time Series Classification via In-Context Learning. *arXiv preprint arXiv:2607.07500* (2026).
 - [36] Bo Lan, Pei Li, Jiaxi Yin, Yunpeng Song, Ge Wang, Han Ding, Jinsong Han, and Fei Wang. 2025. XRF V2: A Dataset for Action Summarization with Wi-Fi Signals, and IMUs in Phones, Watches, Earbuds, and Glasses. *Proceedings of the ACM on Interactive, Mobile, Wearable and Ubiquitous Technologies (IMWUT)* 9, 3 (2025), 1–41. <https://doi.org/10.1145/3749521>
 - [37] Patrick Lewis, Ethan Perez, Aleksandra Piktus, Fabio Petroni, Vladimir Karpukhin, Naman Goyal, Heinrich Küttler, Mike Lewis, Wen-tau Yih, Tim Rocktäschel, Sebastian Riedel, and Douwe Kiela. 2020. Retrieval-Augmented Generation for Knowledge-Intensive NLP Tasks. In *Advances in Neural Information Processing Systems 33 (NeurIPS)*. Curran Associates, Virtual, 9459–9474.
 - [38] Zechen Li, Baiyu Chen, Hao Xue, and Flora D. Salim. 2026. ZARA: Training-Free Motion Time-Series Reasoning via Evidence-Grounded LLM Agents. In *Proceedings of the 64th Annual Meeting of the Association for Computational Linguistics (ACL)*. Association for Computational Linguistics, 14986–15008. <https://doi.org/10.18653/v1/2026.acl-long.684>

arXiv:2508.04038.

- [39] Zechen Li, Shohreh Deldari, Linyao Chen, Hao Xue, and Flora D. Salim. 2025. SensorLLM: Aligning Large Language Models with Motion Sensors for Human Activity Recognition. In *Proceedings of the 2025 Conference on Empirical Methods in Natural Language Processing (EMNLP)*. Association for Computational Linguistics, Suzhou, China, 354–379. <https://doi.org/10.18653/v1/2025.emnlp-main.19>
- [40] Jiayang Liu, Lin Zhong, Jehan Wickramasuriya, and Venu Vasudevan. 2009. uWave: Accelerometer-based personalized gesture recognition and its applications. *Pervasive and Mobile Computing* 5, 6 (2009), 657–675. <https://doi.org/10.1016/j.pmcj.2009.07.007>
- [41] Aleksey Logacjov, Kerstin Bach, Atle Kongsvold, Hilde Bremseth Bårdstu, and Paul Jarle Mork. 2021. HARTH: A Human Activity Recognition Dataset for Machine Learning. *Sensors* 21, 23 (2021), 7853. <https://doi.org/10.3390/s21237853>
- [42] Yunfei Luo, Yuliang Chen, Asif Salekin, and Tauhidur Rahman. 2026. Toward Foundation Model for Multivariate Wearable Sensing of Physiological Signals. *ACM Transactions on Computing for Healthcare* 7, 3 (2026), 1–43. <https://doi.org/10.1145/3803808> Preprint: arXiv:2412.09758.
- [43] Mohammad Malekzadeh, Richard G. Clegg, Andrea Cavallaro, and Hamed Haddadi. 2019. Mobile Sensor Data Anonymization. In *Proceedings of the International Conference on Internet of Things Design and Implementation (IoTDI)*. ACM, Montreal, Canada, 49–58. <https://doi.org/10.1145/3302505.3310068>
- [44] Miguel Matey-Sanz, Sven Casteleyn, and Carlos Granell. 2023. Dataset of Inertial Measurements of Smartphones and Smartwatches for Human Activity Recognition. *Data in Brief* 51 (2023), 109809. <https://doi.org/10.1016/j.dib.2023.109809>
- [45] Conor McCarthy, Loes Quirijnen, Jan Peter van Zandwijk, Zeno Geradts, and Marcel Worring. 2025. Hi-OSCAR: Hierarchical Open-set Classifier for Human Activity Recognition. *Proceedings of the ACM on Interactive, Mobile, Wearable and Ubiquitous Technologies (IMWUT)* 9, 4, Article 199 (2025), 27 pages. <https://doi.org/10.1145/3770681>
- [46] Shenghuan Miao and Ling Chen. 2024. GOAT: A Generalized Cross-Dataset Activity Recognition Framework with Natural Language Supervision. *Proceedings of the ACM on Interactive, Mobile, Wearable and Ubiquitous Technologies (IMWUT)* 8, 4, Article 184 (2024), 28 pages. <https://doi.org/10.1145/3699736>
- [47] Daniela Micucci, Marco Mobilio, and Paolo Napoletano. 2017. UniMiB SHAR: A Dataset for Human Activity Recognition Using Acceleration Data from Smartphones. *Applied Sciences* 7, 10 (2017), 1101. <https://doi.org/10.3390/app7101101>
- [48] Kengo Nakata, Youyang Ng, Daisuke Miyashita, Asuka Maki, Yu-Chieh Lin, and Jun Deguchi. 2022. Revisiting a kNN-Based Image Classification System with High-Capacity Storage. In *Computer Vision – ECCV 2022, 17th European Conference, Proceedings, Part XXXVII (Lecture Notes in Computer Science)*, Vol. 13697. Springer, Tel Aviv, Israel, 457–474. https://doi.org/10.1007/978-3-031-19836-6_26
- [49] Girish Narayanswamy, Xin Liu, Kumar Ayush, Yuzhe Yang, Xu hai Xu, Shun Liao, Jake Garrison, Shyam A. Tailor, Jacob E. Sunshine, Yun Liu, Tim Althoff, Shrikanth Narayanan, Pushmeet Kohli, Jiening Zhan, Mark Malhotra, Shwetak N. Patel, Samy Abdel-Ghaffar, and Daniel McDuff. 2025. Scaling Wearable Foundation Models. In *The Thirteenth International Conference on Learning Representations (ICLR)*. OpenReview.net, Singapore. arXiv:2410.13638.
- [50] Long-Van Nguyen-Dinh, Alberto Calatroni, and Gerhard Tröster. 2014. Robust Online Gesture Recognition with Crowd-sourced Annotations. *Journal of Machine Learning Research (JMLR)* 15, 1 (2014), 3187–3220.
- [51] Yuqi Nie, Nam H. Nguyen, Phanwadee Sinthong, and Jayant Kalagnanam. 2023. A Time Series is Worth 64 Words: Long-term Forecasting with Transformers. In *The Eleventh International Conference on Learning Representations (ICLR)*. OpenReview.net, Kigali, Rwanda.
- [52] Mohammad Norouzi, Tomas Mikolov, Samy Bengio, Yoram Singer, Jonathon Shlens, Andrea Frome, Greg S. Corrado, and Jeffrey Dean. 2014. Zero-Shot Learning by Convex Combination of Semantic Embeddings. In *International Conference on Learning Representations (ICLR)*. arXiv:1312.5650.
- [53] Francisco Javier Ordóñez and Daniel Roggen. 2016. Deep Convolutional and LSTM Recurrent Neural Networks for Multimodal Wearable Activity Recognition. *Sensors* 16, 1 (2016), 115. <https://doi.org/10.3390/s16010115>
- [54] Xiaomin Ouyang, Xian Shuai, Jiayu Zhou, Ivy Wang Shi, Zhiyuan Xie, Guoliang Xing, and Jianwei Huang. 2022. Cosmo: Contrastive Fusion Learning with Small Data for Multimodal Human Activity Recognition. In *Proceedings of the 28th Annual International Conference on Mobile Computing and Networking (MobiCom)*. ACM, 324–337. <https://doi.org/10.1145/3495243.3560519>
- [55] Xiaomin Ouyang, Zhiyuan Xie, Jiayu Zhou, Jianwei Huang, and Guoliang Xing. 2021. ClusterFL: A Similarity-Aware Federated Learning System for Human Activity Recognition. In *Proceedings of the 19th Annual International Conference on Mobile Systems, Applications, and Services (ACM MobiSys)*. ACM, 54–66. <https://doi.org/10.1145/3458864.3467681>
- [56] Alec Radford, Jong Wook Kim, Chris Hallacy, Aditya Ramesh, Gabriel Goh, Sandhini Agarwal, Girish Sastry, Amanda Askell, Pamela Mishkin, Jack Clark, Gretchen Krueger, and Ilya Sutskever. 2021. Learning Transferable Visual Models From Natural Language Supervision. In *Proceedings of the 38th International Conference on Machine Learning (ICML)*. PMLR, Virtual, 8748–8763.
- [57] Nils Reimers and Iryna Gurevych. 2019. Sentence-BERT: Sentence Embeddings using Siamese BERT-Networks. In *Proceedings of the 2019 Conference on Empirical Methods in Natural Language Processing (EMNLP)*. Association for Computational Linguistics, Hong Kong, China, 3982–3992. <https://doi.org/10.18653/v1/D19-1410>
- [58] Attila Reiss and Didier Stricker. 2012. Introducing a New Benchmarked Dataset for Activity Monitoring. In *Proceedings of the 16th International Symposium on Wearable Computers (ISWC)*. IEEE, Newcastle, UK, 108–109. <https://doi.org/10.1109/ISWC.2012.13>
- [59] Jorge-Luis Reyes-Ortiz, Luca Oneto, Albert Samà, Xavier Parra, and Davide Anguita. 2016. Transition-Aware Human Activity Recognition Using Smartphones. *Neurocomputing* 171 (2016), 754–767. <https://doi.org/10.1016/j.neucom.2015.07.085>
- [60] Murat Sensoy, Lance Kaplan, and Melih Kandemir. 2018. Evidential Deep Learning to Quantify Classification Uncertainty. In *Advances in Neural Information Processing Systems 31 (NeurIPS)*. Curran Associates, Montreal, Canada, 3183–3193.
- [61] Jeff Shen, Francois Lanusse, Liam Holden Parker, Ollie Liu, Tom Hehir, Leopoldo Sarra, Lucas Meyer, Micah Bowles, Sebastian Wagner-Carena, Helen Qu, Siavash Golkar, Alberto Bietti, Kyunghyun Cho, Miles Cranmer, and Shirley Ho. 2025.

- Universal Spectral Tokenization via Self-Supervised Panchromatic Representation Learning. In *NeurIPS 2025 Workshop on Machine Learning and the Physical Sciences (ML4PS)*. San Diego, CA. arXiv:2510.17959.
- [62] Muhammad Shoaib, Stephan Bosch, Ozlem Durmaz Incel, Hans Scholten, and Paul J. M. Havinga. 2014. Fusion of Smartphone Motion Sensors for Physical Activity Recognition. *Sensors* 14, 6 (2014), 10146–10176. <https://doi.org/10.3390/s140610146>
- [63] Muhammad Shoaib, Stephan Bosch, Ozlem Durmaz Incel, Hans Scholten, and Paul J. M. Havinga. 2016. Complex Human Activity Recognition Using Smartphone and Wrist-Worn Motion Sensors. *Sensors* 16, 4 (2016), 426. <https://doi.org/10.3390/s16040426>
- [64] Luis Sigcha, Carlos Polvorinos-Fernández, Néelson Costa, Susana Costa, Pedro Arezes, Miguel Gago, Chaiwoo Lee, Juan Manuel López, Guillermo de Arcas, and Ignacio Pavón. 2023. Monipar: Movement Data Collection Tool to Monitor Motor Symptoms in Parkinson’s Disease Using Smartwatches and Smartphones. *Frontiers in Neurology* 14 (2023), 1326640. <https://doi.org/10.3389/fneur.2023.1326640>
- [65] Niloy Sikder and Abdullah-Al Nahid. 2021. KU-HAR: An Open Dataset for Heterogeneous Human Activity Recognition. *Pattern Recognition Letters* 146 (2021), 46–54. <https://doi.org/10.1016/j.patrec.2021.02.024>
- [66] Nirhoshan Sivaroopan, Hansi Karunarathna, Chamara Madarasingha, Anura Jayasumana, and Kanchana Thilakarathna. 2026. RAG-HAR: Retrieval Augmented Generation-based Human Activity Recognition. In *Proceedings of the IEEE International Conference on Pervasive Computing and Communications (PerCom)*. IEEE. <https://doi.org/10.1109/PERCOM67906.2026.11524560> arXiv:2512.08984.
- [67] Jake Snell, Kevin Swersky, and Richard Zemel. 2017. Prototypical Networks for Few-shot Learning. In *Advances in Neural Information Processing Systems 30 (NIPS)*. Curran Associates, Long Beach, CA, 4077–4087.
- [68] Allan Stisen, Henrik Blunck, Sourav Bhattacharya, Thor Siger Prentow, Mikkel Bent Kjærgaard, Anind Dey, Tobias Sonne, and Mads Møller Jensen. 2015. Smart Devices Are Different: Assessing and Mitigating Mobile Sensing Heterogeneities for Activity Recognition. In *Proceedings of the 13th ACM Conference on Embedded Networked Sensor Systems (SenSys)*. ACM, Seoul, South Korea, 127–140.
- [69] David Strömbäck, Sangxia Huang, and Valentin Radu. 2020. MM-Fit: Multimodal Deep Learning for Automatic Exercise Logging across Sensing Devices. *Proceedings of the ACM on Interactive, Mobile, Wearable and Ubiquitous Technologies (IMWUT)* 4, 4 (2020), 1–22. <https://doi.org/10.1145/3432701>
- [70] Jianlin Su, Murtadha Ahmed, Yu Lu, Shengfeng Pan, Wen Bo, and Yunfeng Liu. 2024. RoFormer: Enhanced Transformer with Rotary Position Embedding. *Neurocomputing* 568 (2024), 127063. <https://doi.org/10.1016/j.neucom.2023.127063>
- [71] Timo Sztyler and Heiner Stuckenschmidt. 2016. On-Body Localization of Wearable Devices: An Investigation of Position-Aware Activity Recognition. In *Proceedings of the IEEE International Conference on Pervasive Computing and Communications (PerCom)*. IEEE, Sydney, Australia, 1–9. <https://doi.org/10.1109/PERCOM.2016.7456521>
- [72] Aaron van den Oord, Yazhe Li, and Oriol Vinyals. 2018. Representation Learning with Contrastive Predictive Coding. *arXiv preprint arXiv:1807.03748* (2018).
- [73] Ashish Vaswani, Noam Shazeer, Niki Parmar, Jakob Uszkoreit, Llion Jones, Aidan N. Gomez, Lukasz Kaiser, and Illia Polosukhin. 2017. Attention Is All You Need. In *Advances in Neural Information Processing Systems 30 (NeurIPS)*. Curran Associates, Long Beach, CA, 5998–6008.
- [74] George Vavoulas, Charikleia Chatzaki, Thodoris Malliotakis, Matthew Padiaditis, and Manolis Tsiknakis. 2016. The Mobi-Act Dataset: Recognition of Activities of Daily Living Using Smartphones. In *Proceedings of the International Conference on Information and Communication Technologies for Ageing Well and e-Health (ICT4AWE)*. SciTePress, Rome, Italy, 143–151. <https://doi.org/10.5220/0005792401430151>
- [75] Oriol Vinyals, Charles Blundell, Timothy Lillicrap, Koray Kavukcuoglu, and Daan Wierstra. 2016. Matching Networks for One Shot Learning. In *Advances in Neural Information Processing Systems 29 (NIPS)*. Curran Associates, Barcelona, Spain, 3630–3638.
- [76] Qingxin Wei, Jiaming Huang, Yi Gao, and Wei Dong. 2025. One Model to Fit Them All: Universal IMU-based Human Activity Recognition with LLM-assisted Cross-dataset Representation. *Proceedings of the ACM on Interactive, Mobile, Wearable and Ubiquitous Technologies (IMWUT)* 9, 3, Article 139 (2025), 22 pages. <https://doi.org/10.1145/3749509>
- [77] Gary M. Weiss, Kenichi Yoneda, and Thaier Hayajneh. 2019. Smartphone and Smartwatch-Based Biometrics Using Activities of Daily Living. *IEEE Access* 7 (2019), 133190–133202. <https://doi.org/10.1109/ACCESS.2019.2940729>
- [78] Yongqin Xian, Christoph H. Lampert, Bernt Schiele, and Zeynep Akata. 2019. Zero-Shot Learning—A Comprehensive Evaluation of the Good, the Bad and the Ugly. *IEEE Transactions on Pattern Analysis and Machine Intelligence* 41, 9 (2019), 2251–2265. <https://doi.org/10.1109/TPAMI.2018.2857768>
- [79] Huatao Xu, Pengfei Zhou, Rui Tan, and Mo Li. 2023. Practically Adopting Human Activity Recognition. In *Proceedings of the 29th Annual International Conference on Mobile Computing and Networking (ACM MobiCom)*. ACM, Article 85, 15 pages. <https://doi.org/10.1145/3570361.3613299>
- [80] Huatao Xu, Pengfei Zhou, Rui Tan, Mo Li, and Guobin Shen. 2021. LIMU-BERT: Unleashing the Potential of Unlabeled Data for IMU Sensing Applications. In *Proceedings of the 19th ACM Conference on Embedded Networked Sensor Systems (SenSys)*. ACM, Coimbra, Portugal, 220–233. <https://doi.org/10.1145/3485730.3485937>
- [81] Maxwell A. Xu, Jaya Narain, Gregory Darnell, Haraldu Tomás Hallgrímsson, Hyewon Jeong, Darren Forde, Richard Fineman, Karthik Jayaraman Raghuram, James M. Rehg, and Shirley You Ren. 2025. RelCon: Relative Contrastive Learning for a Motion Foundation Model for Wearable Data. In *The Thirteenth International Conference on Learning Representations (ICLR)*. OpenReview.net, Singapore. arXiv:2411.18822.
- [82] Maxwell A. Xu, Girish Narayanswamy, Kumar Ayush, Dimitris Spathis, Shun Liao, Shyam A. Tailor, Ahmed Metwally, A. Ali Heydari, Yuwei Zhang, Jake Garrison, Samy Abdel-Ghaffar, Xuhai Xu, Ken Gu, Jacob E. Sunshine, Ming-Zher Poh, Yun Liu, Tim Althoff, Shrikanth Narayanan, Pushmeet Kohli, Mark Malhotra, Shwetank N. Patel, Yuzhe Yang, James M. Rehg, Xin Liu, and Daniel McDuff. 2025. LSM-2: Learning from Incomplete Wearable Sensor Data. *arXiv preprint arXiv:2506.05321* (2025). <https://doi.org/10.48550/arXiv.2506.05321>

- [83] Zongzhe Xu, Aakarsh Anand, Sarah Jiang, Chuntung Zhuang, Zitao Shuai, Sriram Sankararaman, and Yuzhe Yang. 2026. Inertia-1: An Open Exploration of Wearable Motion Foundation Models. *arXiv preprint arXiv:2607.06617* (2026).
- [84] Hua Yan, Heng Tan, Yi Ding, Pengfei Zhou, Vinod Nambodiri, and Yu Yang. 2025. Large Language Model-guided Semantic Alignment for Human Activity Recognition. *Proceedings of the ACM on Interactive, Mobile, Wearable and Ubiquitous Technologies (IMWUT)* 9, 4 (2025), 1–25. <https://doi.org/10.1145/3770652> System name: LanHAR.
- [85] Yan Yan, Dali Chen, Yushi Liu, Jinjin Zhao, Bo Wang, Xuankun Wu, Xiaohao Jiao, Yuqian Chen, Huihui Li, and Xuchao Ren. 2021. TNDA-HAR. *IEEE DataPort*. (2021). <https://doi.org/10.21227/4epb-pg26>
- [86] Hang Yuan, Shing Chan, Andrew P. Creagh, Catherine Tong, Aidan Acquah, David A. Clifton, and Aiden Doherty. 2024. Self-supervised learning for human activity recognition using 700,000 person-days of wearable data. *npj Digital Medicine* 7, 1, Article 91 (2024), 10 pages. <https://doi.org/10.1038/s41746-024-01062-3> arXiv:2206.02909.
- [87] Hao Zhang, Zhan Zhuang, Xuehao Wang, Xiaodong Yang, and Yu Zhang. 2025. MoPFormer: Motion-Primitive Transformer for Wearable-Sensor Activity Recognition. In *Advances in Neural Information Processing Systems 38 (NeurIPS)*. Curran Associates, San Diego, CA. <https://doi.org/10.52202/085713-1313> arXiv:2505.20744.
- [88] Mi Zhang and Alexander A. Sawchuk. 2012. USC-HAD: A Daily Activity Dataset for Ubiquitous Activity Recognition Using Wearable Sensors. In *Proceedings of the 2012 ACM Conference on Ubiquitous Computing (UbiComp)*. ACM, Pittsburgh, PA, 1036–1043. <https://doi.org/10.1145/2370216.2370438>
- [89] Xiyuan Zhang, Danielle C. Maddix, Junming Yin, Nick Erickson, Abdul Fatir Ansari, Boran Han, Shuai Zhang, Leman Akoglu, Christos Faloutsos, Michael W. Mahoney, Cuixiong Hu, Huzefa Rangwala, George Karypis, and Bernie Wang. 2025. Mitra: Mixed Synthetic Priors for Enhancing Tabular Foundation Models. In *Advances in Neural Information Processing Systems (NeurIPS)*. arXiv:2510.21204.
- [90] Xiyuan Zhang, Diyan Teng, Ranak Roy Chowdhury, Shuheng Li, Dezhi Hong, Rajesh K. Gupta, and Jingbo Shang. 2024. UniMTS: Unified Pre-training for Motion Time Series. In *Advances in Neural Information Processing Systems 37 (NeurIPS)*. <https://doi.org/10.52202/079017-3414>
- [91] Yuwei Zhang, Kumar Ayush, Siyuan Qiao, A. Ali Heydari, Girish Narayanswamy, Maxwell A. Xu, Ahmed A. Metwally, Shawn Xu, Jake Garrison, Xuhai Xu, Tim Althoff, Yun Liu, Pushmeet Kohli, Jiening Zhan, Mark Malhotra, Shwetak Patel, Cecilia Mascolo, Xin Liu, Daniel McDuff, and Yuzhe Yang. 2025. SensorLM: Learning the Language of Wearable Sensors. In *Advances in Neural Information Processing Systems 38 (NeurIPS)*. Curran Associates, San Diego, CA. <https://doi.org/10.52202/085713-1559> arXiv:2506.09108.
- [92] Shibo Zhao, Sifan Zhou, Raphael Blanchard, Yuheng Qiu, Wenshan Wang, and Sebastian Scherer. 2025. Tartan IMU: A Light Foundation Model for Inertial Positioning in Robotics. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*. <https://doi.org/10.1109/CVPR52734.2025.02097>

A SCOPE OF THE CLOSEST SYSTEMS

Table 5 keeps four questions separate: handling acquisition rate, receiving placement or channel context, accepting run-time activity names, and evaluating multiple target example counts without parameter updates. A zero-shot result is not marked as a k curve, and a linear probe or fine-tuning result is not marked as no-update enrollment.

B CORPUS CONSTRUCTION

Table 3 is generated from each stream’s `grids/native/*meta.json`. Window counts are properties of the HALO converters and grid policy, not quantities claimed by the source dataset papers. The evidence ledger records the aggregation command and distinguishes materialized rows, quality-screened rows, subject splits, and the balanced 250,000-window bank.




Figure 4: Materialized configurations by stored grid rate. The current roster contains 56 Phase-A configurations and 14 external-evaluation configurations. Source acquisition rate is stored separately and bounds observability when a converter has already resampled a stream.

The expanded training roster deliberately includes body-worn stress configurations from DSADS and Opportunity; an earlier draft incorrectly listed both as excluded. Six newly added training sources contribute most of the placement coverage: DSADS, FORTH-TRACE, Opportunity, RealDISP, MM-Fit, and PhyTMO. MoniPar, SPAR, and Upper Limb Use are held out.

Sources not used in the named recipe. HAPT [59] reuses the UCI-HAR participants and recordings, so it is excluded to avoid near-duplicate subjects. RecGym [9] distributes per-axis min-max-normalized data that no longer preserve the physical scale used by the tokenizer.

Table 5: Reported scope of related systems. Entries summarize the cited evaluation rather than infer unreported capabilities. “Task fit” includes a probe or fine-tuning; “zero only” is not a target-example curve.

System	Rate/input treatment	Configuration context	Label interface	Target adaptation
LiMU-BERT [80]	common 20 Hz	none	fixed downstream	task fit
CrossHAR [27]	common-rate inputs	none	fixed downstream	task fit
harnet [86]	wrist contract	wrist accelerometer	fixed downstream	task fit
LSM / LSM-2 [49, 82]	multimodal, incomplete data	modality availability	fixed tasks	probe/fine-tune
UniMTS [90]	resampled inputs	placement graph	text-conditioned tasks	task fit
NormWear [42]	common-rate input	variable channels	text interface	zero/task evaluation
ZeroHAR [17]	configured input	categorical context	run-time names	zero only
CHARM / SensorsVoice [19, 20]	channel-tokenized	channel descriptions	task dependent	task evaluation
MoPFormer [87]	all data at 100 Hz	wearable context	fixed tasks	task fit
FlowState [25]	rate equivariant	sampling rate	forecasting output	not HAR
RAG-HAR [66]	retrieved descriptors	sensor evidence	run-time names	zero only
ZARA [38]	multi-sensor evidence	sensor evidence	run-time names	zero only
SensorLM [91]	wearable pretraining	sensor/task prompts	task dependent	task adaptation
HALO	stored-grid bands in Hz	per-sensor text + measurements	run-time names	$k = 0, 1, 2, 4, 8$, no update

Table 6: Planned baseline groups. Results are not reported until every row has an artifact using the same support and query cohorts.

Model	Weight source	Input contract	Candidate bridge
LiMU-BERT [80]	corpus matched	20 Hz, padded 6-channel	ConSE
CrossHAR [27]	corpus matched	common rate, padded channels	ConSE
harnet [86]	released	30 Hz wrist acceleration	ConSE
UniMTS [90]	released	placement graph	native interface
NormWear [42]	released	model-specific common rate	text interface
HALO	current run	per-sensor stored-grid input	evidence rule

The MobiAct [74] download did not produce a materialized grid and is not counted. HARTH [41] is available only as an optional placement stress source, and ExtraSensory, NHANES, HMOG, and KneE-PAD are optional scale/stress sources rather than members of the reported 18-dataset run.

C BASELINE CONTRACTS

Models receive different input information, so a future comparison must disclose both weight source and input contract. LiMU-BERT and CrossHAR require corpus-matched retraining. Released harnet, UniMTS, and NormWear weights were trained on different corpora. ConSE [52] can translate a fixed training classifier into candidate text scores, but it does not make the underlying encoder open label or provide an enrollment mechanism.

D TOKENIZER CHECKS REQUIRED FOR PUBLICATION

The implementation contains unit guards for DFT overflow, acquisition-rate observability, sensor folding, and axis validity. Those software tests do not by themselves establish empirical rate-equivalence or posture utility. The retired figures used hard-coded values without an archived generator. Their replacements must be produced by a checked-in script that records input windows, anti-alias filtering, source and stored rates, checkpoint hash, random seed, and raw feature arrays. The signed-low-frequency ablation must use the same evaluation cohort and report uncertainty.

E PHASE-A SPECIFICATION

The reported run uses a fixed filterbank, sensor-granularity tokens, width 256, six layers, eight heads, batch size 256, 30,000 steps, learning rate 3×10^{-4} , 1,000 warmup steps, weight decay 0.05, gradient clipping at 1.0, and seed 20260718. The two top-level loss coefficients are fixed after calibration at step 2,000: 3.998090 for the JEPa family and 0.689099 for VICReg. Descriptor retrieval has coefficient 0.5 inside the JEPa family. The teacher decay is 0.996.

Short patch durations are drawn from 0.4–0.8 s, long durations from 1.0–1.5 s, with a minimum ratio of 1.75. A dataset with n windows is sampled proportional to $n^{0.25}$ with a maximum share of 25%; subjects are sampled proportional to the square root of their window count. Activity labels do not enter the training sampler or loss. The selected checkpoint is step 4,000, which maximizes the predefined subject-disjoint label-stream-macro validation score over the completed run.

F EVIDENCE ARTIFACT AND CONTROLS

The active bank schema has 250,000 source windows, 56 configurations, and 161 labels. It records 3,942,901 valid patch rows and 7,653,242 per-sensor rows. The gate artifact is bound to this bank, checkpoint, text-embedding path, and evaluation schema. Evaluation refuses legacy banks or foreign artifacts unless the mismatch is explicitly acknowledged.

The prediction budget is 64 evidence rows per query patch and candidate. Sensor-bias similarity is disabled. For semantically unfamiliar sensor or concept text, the gate blends toward a neutral value. The following telemetry is mandatory whenever the gate is reported:

- gate mean, spread, extrema, and vetoed-row fraction;
- same-dataset retrieval with the acquisition-bias blend both off and on;
- held-out concept, dataset, and body-region errors relative to constant and per-concept baselines; and
- within-concept correlation between gate latents and the train-only resolvability table.

Every enrollment result shares one deterministic query cohort and nested support prefixes across the gate, all-one admissibility, support removal, support-label shuffling, prototype, and ridge controls. These controls are part of the claim, not optional ablations.

G EVIDENCE LEDGER

The paper source includes `EVIDENCE_LEDGER.md`. Each quantitative manuscript statement is classified as verified, development-only, sealed, or pending. A verified entry names the producing command or log, source-data manifest, seed, checkpoint, and artifact fingerprint. Values without that chain are not printed as results. This policy is also why the earlier geometry, gravity, rate-equivalence, DFT-cosine, deployment, and phone-latency numbers are absent from this revision.